Jiangjie (Becket) Qin LinkedIn

Mission Critical Messaging with Apache Kafka



促进软件开发领域知识与创新的传播



关注InfoQ官方信息

及时获取QCon软件开发者 大会演讲视频信息



[北京站] 2016年12月2日-3日

咨询热线: 010-89880682



[北京站] 2017年4月16日-18日

咨询热线: 010-64738142

Agenda

- Introduction to Apache Kafka
- Kafka based replication in Espresso
- Message Integrity guarantees
- Performance
- Large message handling
- Security
- Q&A

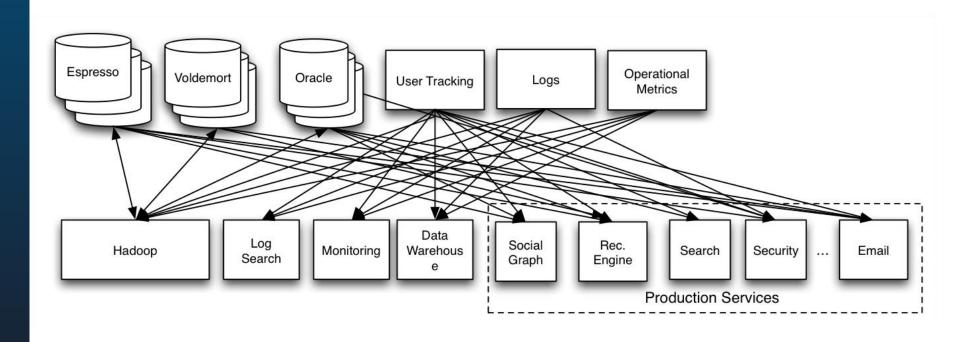


Agenda

- Introduction to Apache Kafka
- Kafka based replication in Espresso
- Message Integrity guarantees
- Performance
- Large message handling
- Security
- Q&A

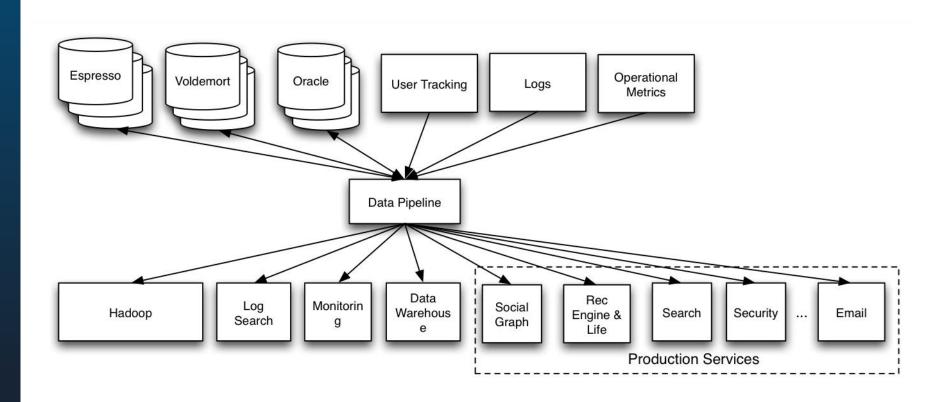


One way to aggregate data



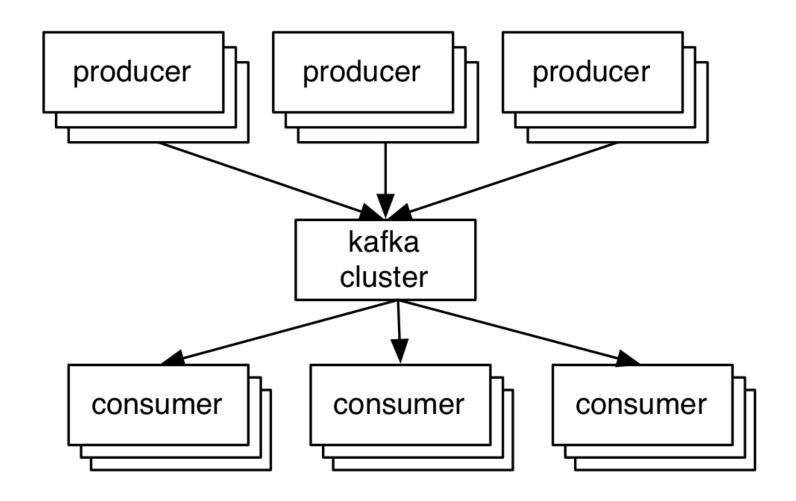


Another way to aggregate data

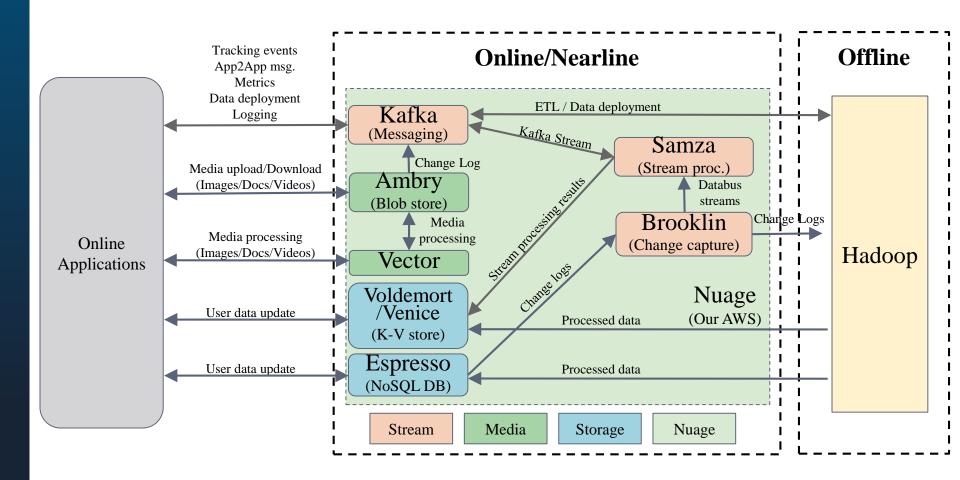




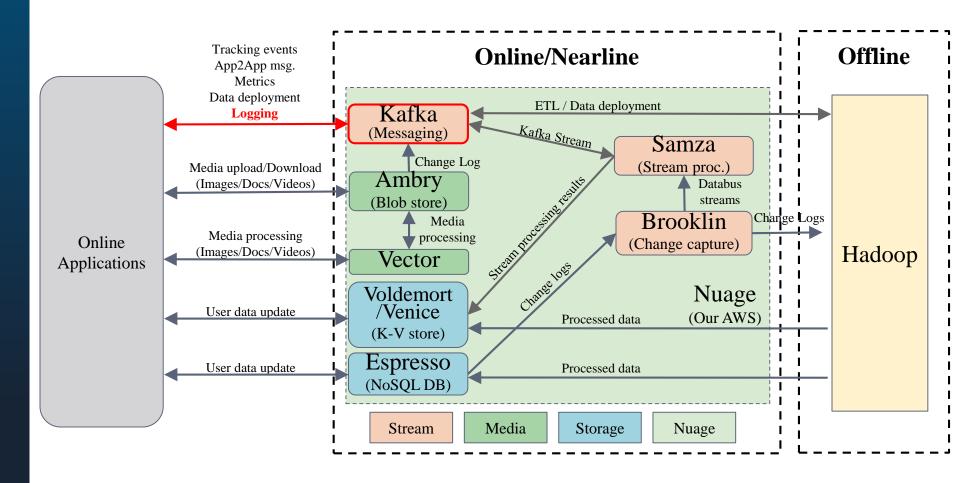
Apache Kafka



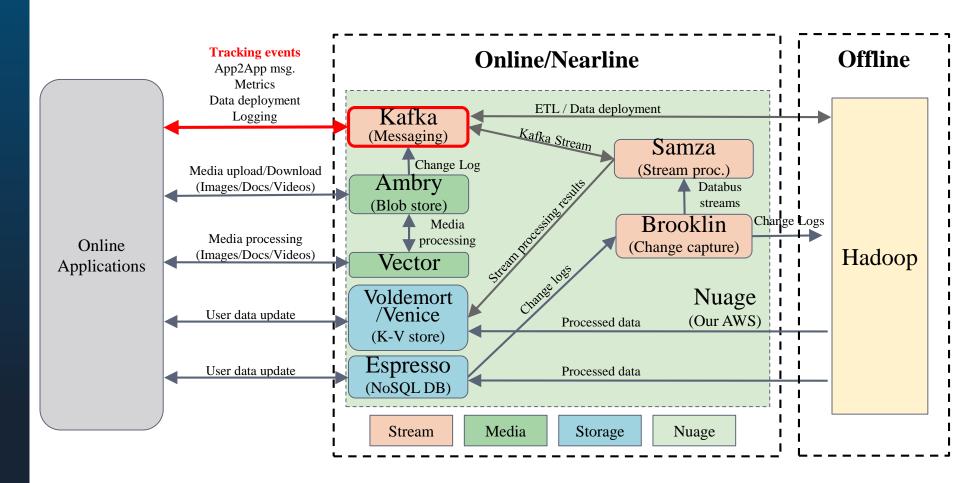




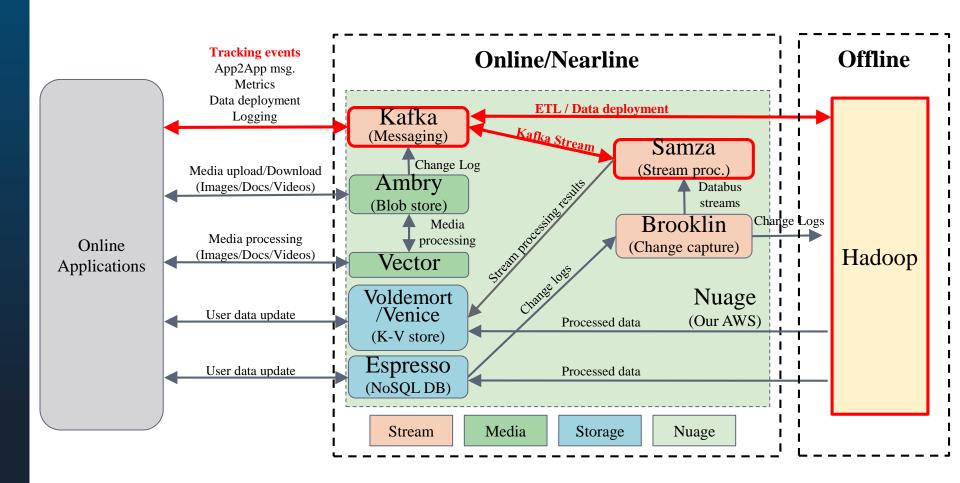




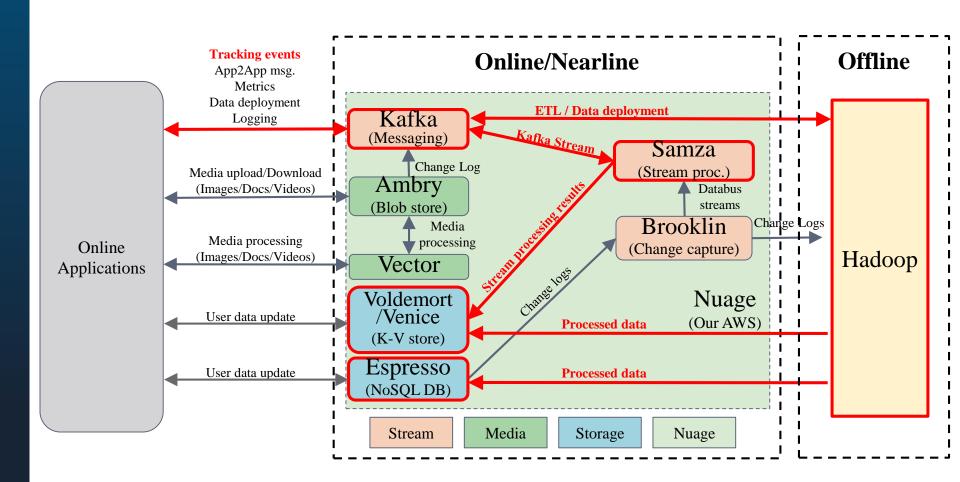
下午16:40, 百宴厅1, LinkedIn基于Kafka和ElasticSearch的实时日志分析



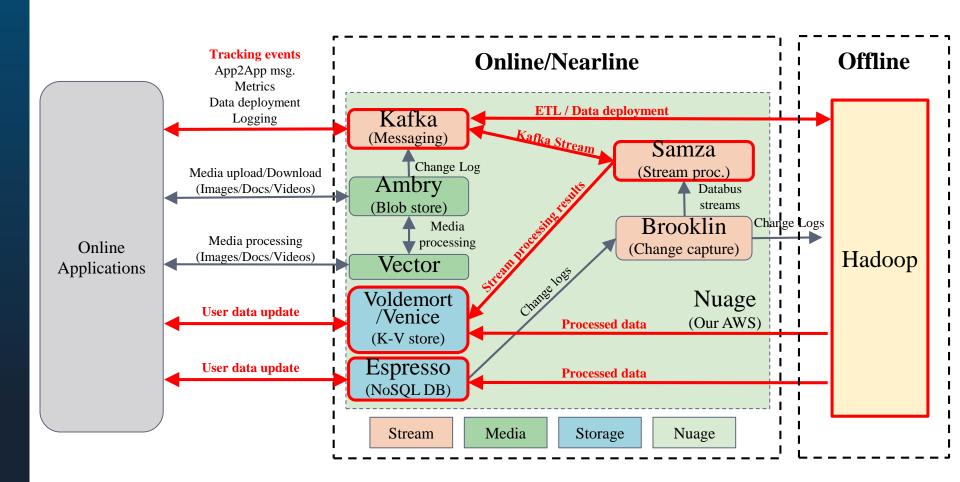




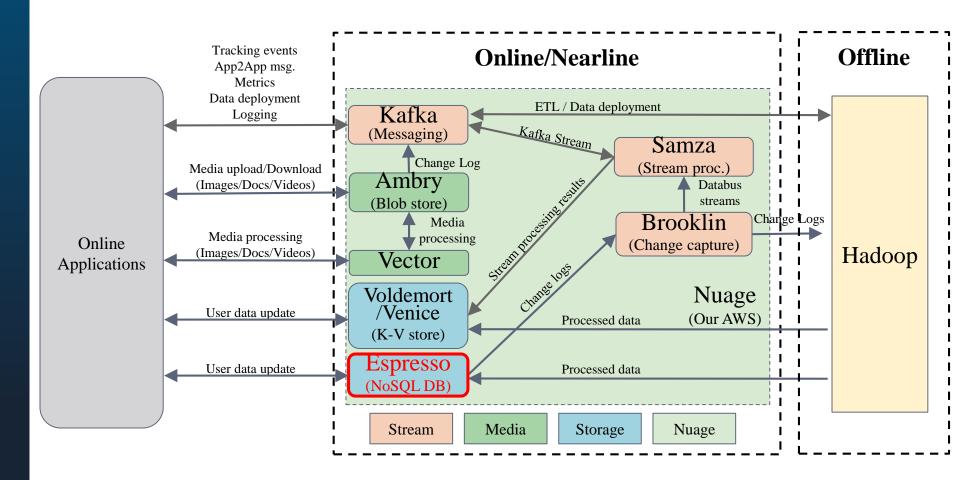












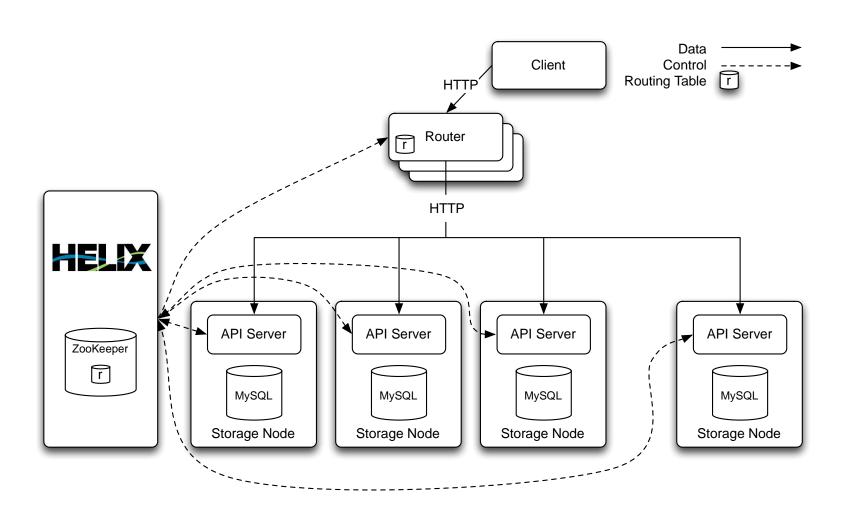


Agenda

- Introduction to Apache Kafka
- Kafka based replication in Espresso
- Message Integrity guarantees
- Performance
- Large message handling
- Security



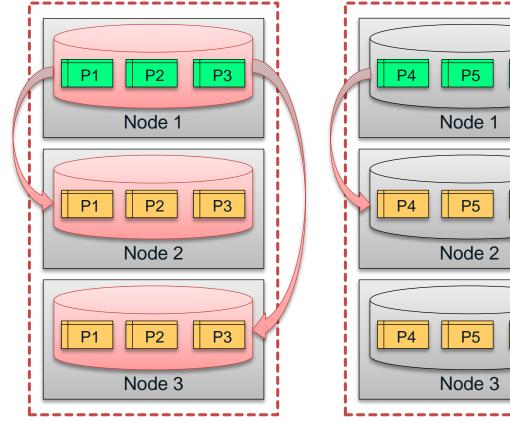
Espresso - A scalable NoSQL DB

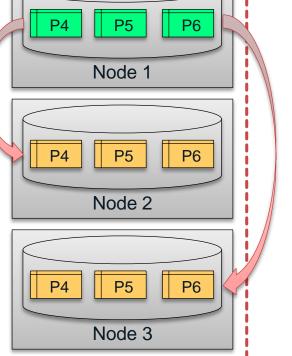


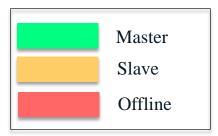


MySQL Based Replication

MySQL instance level replication



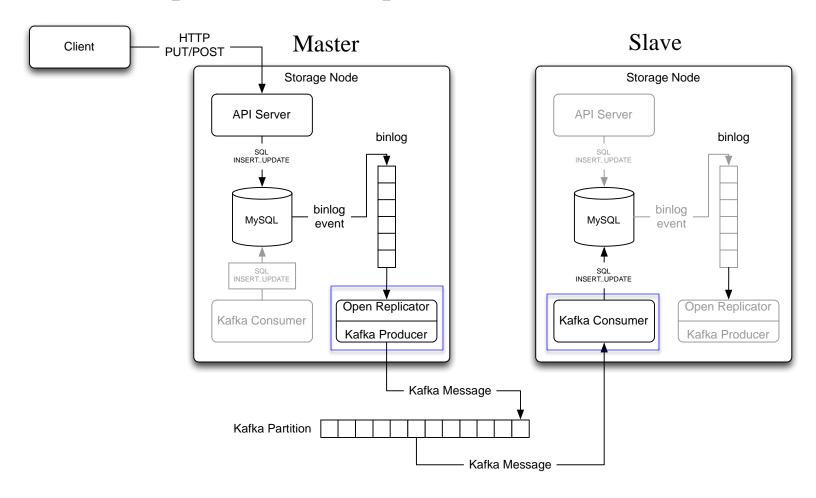






Kafka Based Replication

Kafka based partition level replication





Mission Critical Messaging

- No message loss
- In-order delivery
- Exactly once semantic _
- High throughput
- Low latency
- Handle large messages
- Security







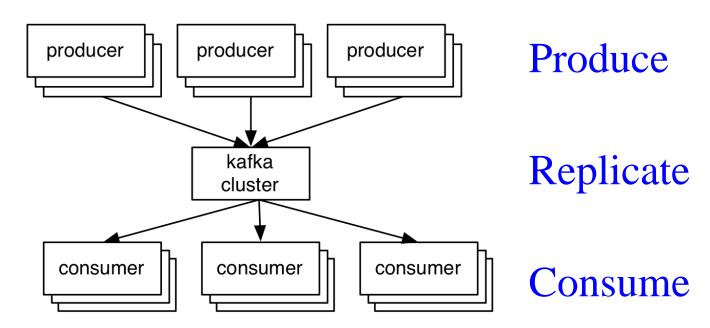
Agenda

- Introduction to Apache Kafka
- Kafka based replication in Espresso
- Message Integrity guarantees
- Performance
- Large message handling
- Security



Data Integrity

- No message loss
- In-order delivery
- Exactly once semantic





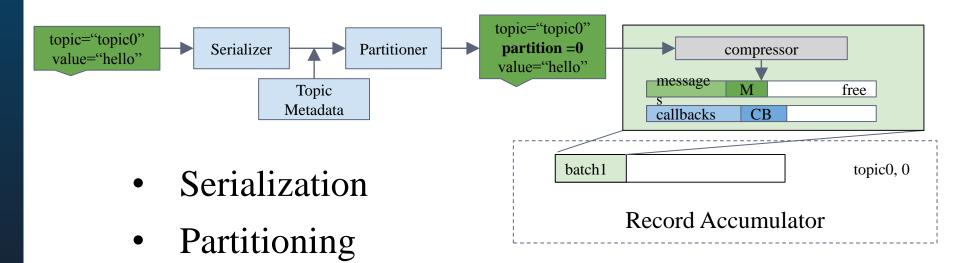
Produce

- A well implemented producer usually supports:
 - Batching the messages
 - Sending the messages asynchronously
- Example:

org.apache.kafka.clients.producer.KafkaProducer



KafkaProducer



- Compression
- ✓ Tasks done by the user thread



KafkaProducer

Sender:

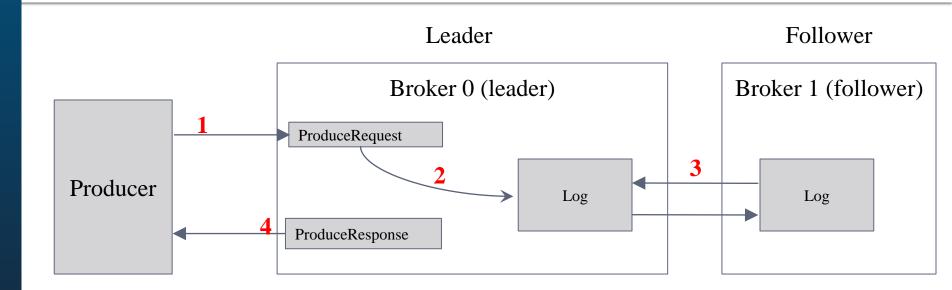
- **1. polls** batches from the batch queues (one batch / partition)
- **2. groups** batches based on the leader broker
- 3. sends the grouped batches to the brokers
- 4. Fire callbacks after receiving the response compressor message free callbacks CB batch1 topic0, 0 Broker 0 callbacks **CB** batch0 topic0, 1 resp Broker 1 batch2 batch0 batch1 topic1, 0 resp sender thread Record Accumulator

Ensure No Message Loss

- Enable retries in the sender
 - (e.g. retries=5)
- Keep the messages not acked yet
 - Espresso only checkpoints at the transaction boundary after the callback is fired
- acks=all



acks=all



- 1. [Network] Send ProduceRequest
- 2. [Broker] Append messages to the leader's log
- 3. [Broker] Replication (before sending the response)
- 4. [Broker] ProduceResponse



Durability guarantee

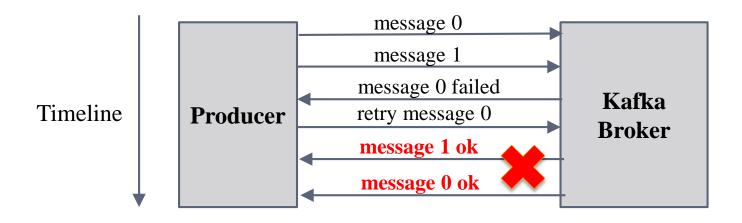
- In-Sync-Replica (ISR)
 - A replica that can keep up with the leader
- Semantic for acks

acks	Throughput	Latency	Durability
no ack(0)	high	low	No guarantee
Leader only(1)	medium	medium	leader
All ISR(-1)	low	high	All ISR



In order delivery

- max.in.flight.requests.per.connection = 1
 - Request pipelining

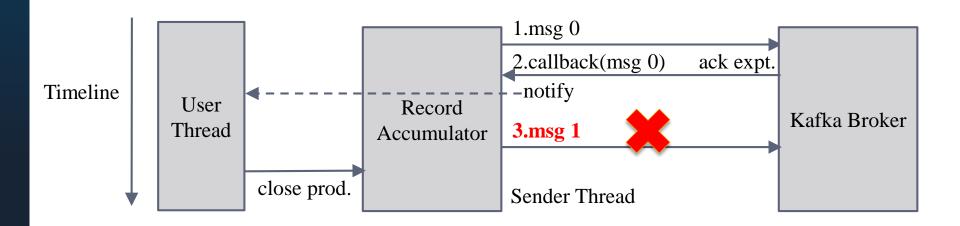


max.in.flight.requests.per.connection=2



In order delivery

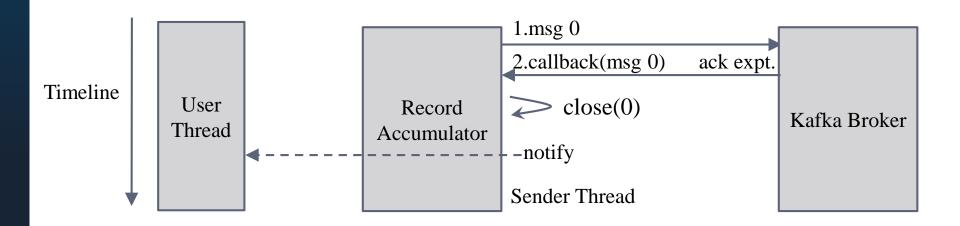
- Close the producer in the callback with 0 timeout on failures
 - Callbacks are executed in the Sender thread





In order delivery

 Close the producer in the callback with 0 timeout on failures





Broker

- min.isr=2
 - If acks=all, at least 2 copies of messages are required on the broker
 - replication.factor needs be 3 to tolerate single broker failure



Replica =
$$\{0, 1\}$$

ISR = $\{0, 1\}$

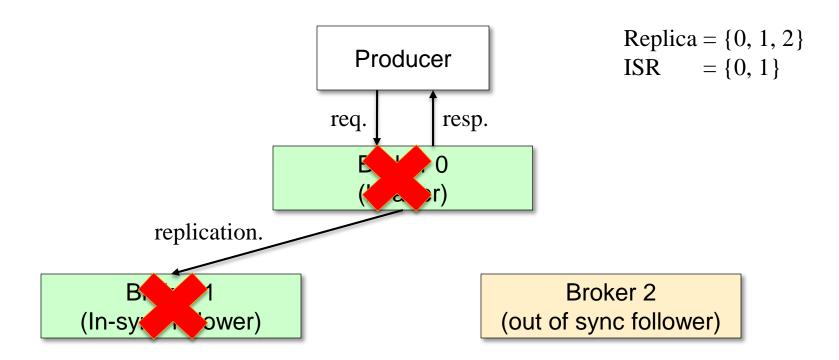
Replica = $\{0, 1\}$

ISR = $\{0\}$



Broker

- unclean.leader.election.enable= false
 - Only in-sync replicas can become leader





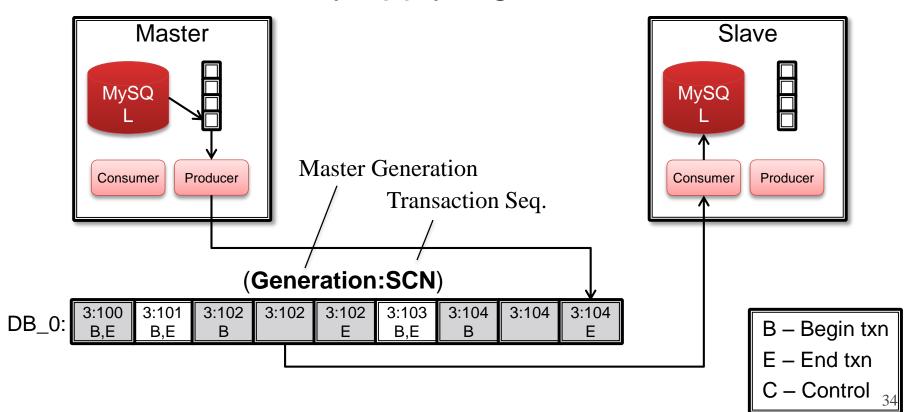
Consumer

- Disable auto offset commit
- Manually commit offset
 - only commit offsets after successfully processing the messages



Consumer

- Exactly once delivery (Espresso)
 - Consumer only apply higher Generation:SCN



Agenda

- Introduction to Apache Kafka
- Kafka based replication in Espresso
- Message Integrity guarantees
- Performance
- Large message handling
- Security

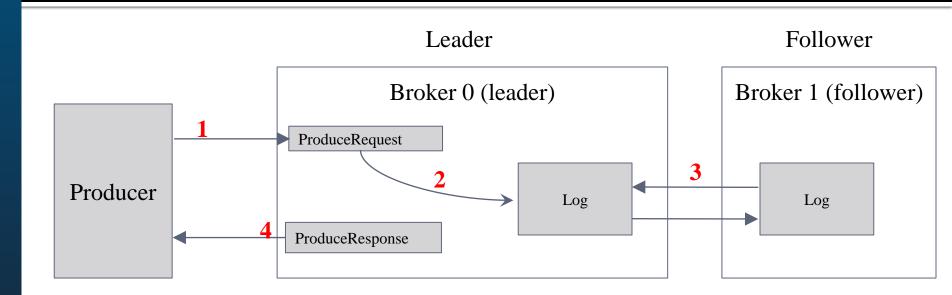


Performance Tuning

- Performance tuning is case by case
 - Traffic pattern sensitive
 - Application requirement specific
- Producer performance is more interesting
 - Especially for acks=all
 - See more
 - Producer performance tuning for Apache Kafka (http://www.slideshare.net/JiangjieQin/producer-performance-tuning-for-apache-kafka-63147600)



Latency when acks=all



- 1. [Network] Send ProduceRequest
- 2. [Broker] Append messages to the leader's log
- 3. [Broker] Replication (synchronously) ———————————————increases latency
- 4. [Broker] ProduceResponse



Latency when acks=all

- Kafka replication is a pull model
- Increase num.replica.fetchers
 - Parallel fetch
- Not perfect solution
 - Diminishing effect (1/N)
 - Scalability concern
 - Replica fetchers per broker = (Cluster_Size 1) * num.replica.fetchers



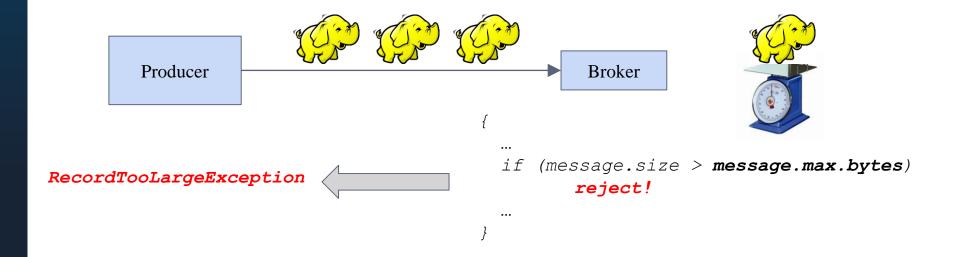
Agenda

- Introduction to Apache Kafka
- Kafka based replication in Espresso
- Message Integrity guarantees
- Performance
- Large message handling
- Security



What is a large message

- Kafka has a limit on the maximum size of a single message
 - Enforced on the compressed wrapper message if compression is used





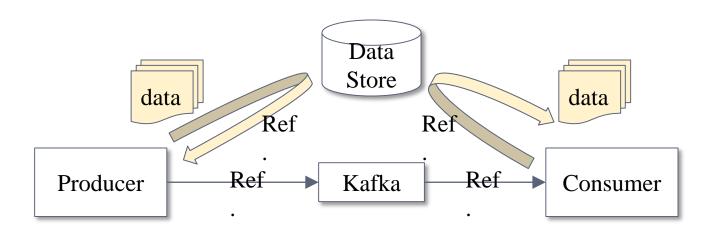
Why limit message size

- Increase the memory pressure in the broker
- Large messages are expensive to handle and could slow down the brokers.
- A reasonable message size limit can handle vast majority of the use cases.



Typical solution

Reference based messaging

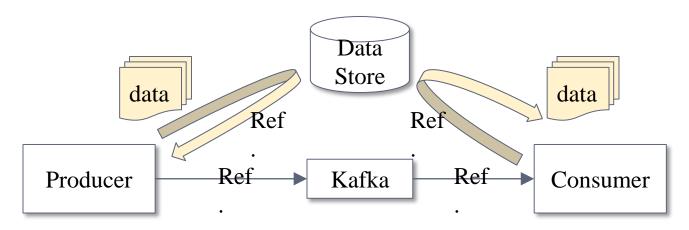




What works

- Unknown maximum row size
- Strict no data loss
- Strict message order guarantee

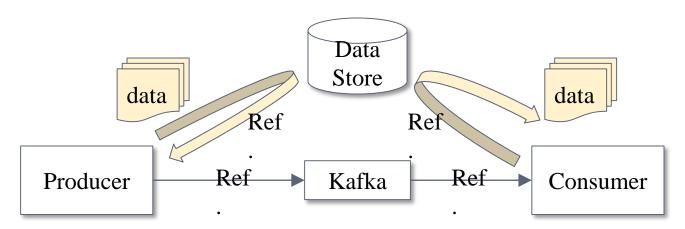
Works fine as long as the durability of the data store can be guaranteed.





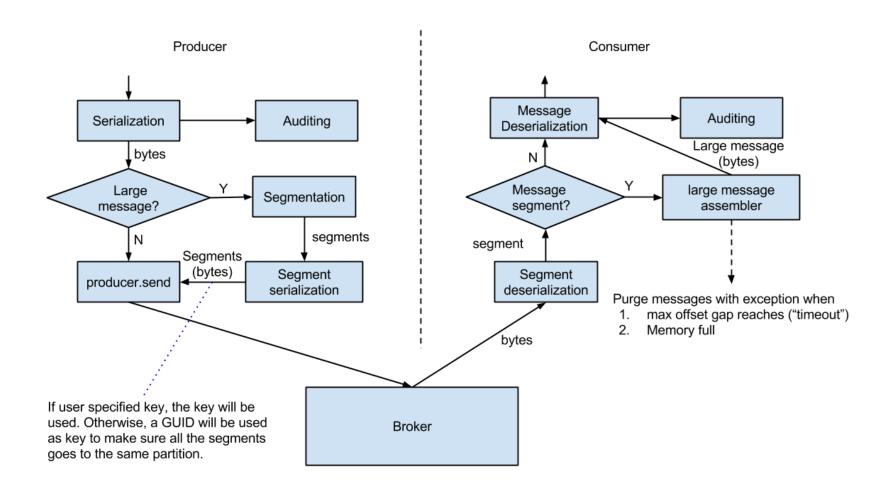
What does not work

- Replicates a data store by using another data store....
- Sporadic large messages
- Low end to end latency
 - There are more round trips in the system.
 - Need to make sure the data store is fast





In-line large message handling



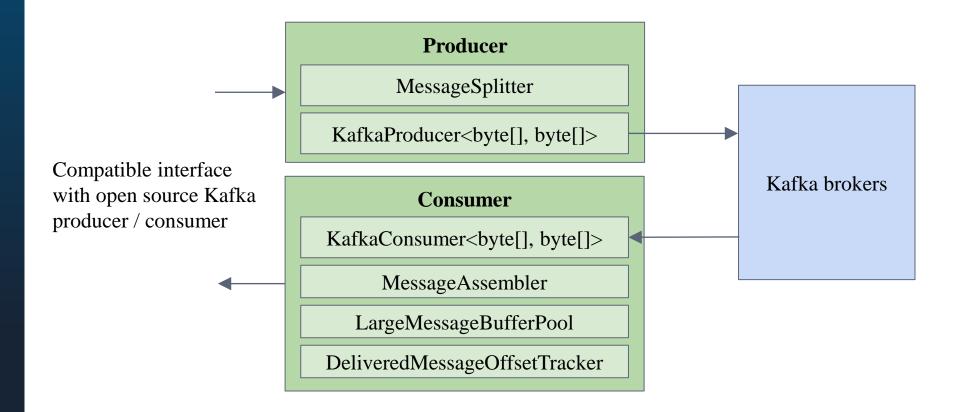


In-line large message support

	Reference Based Messaging	In-line large message support
Operational complexity	Two systems to maintain	Only maintain Kafka
System stability	Depend on : The consistency between Kafka and the external storage The durability of external storage	Only depend on Kafka
Cost to serve	Kafka + External Storage	Only maintain Kafka
End to end latency	Depend on the external storage	The latency of Kafka
Client complexity	Need to deal with envelopes	Much more involved
Functional limitations	Almost none	Some limitations



In-line large message support





Some details

- Many interesting details
 - The offset of a large message
 - Offset Tracking
 - Rebalance and duplicates handling
 - Producer callback
 - Memory management
 - Performance overhead
 - Compatibility with existing messages



Some details

- Many interesting details
 - The offset of a large message
 - Offset Tracking
 - Rebalance and duplicates handling
 - Producer callback
 - Memory management
 - Performance overhead
 - Compatibility with existing messages



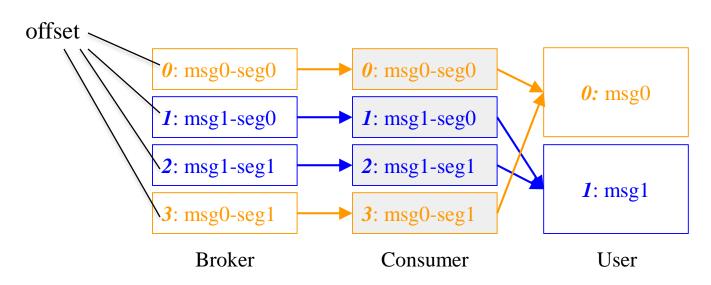
Offset of a large message

- Each message in Kafka has an Offset
 - The logical sequence in the log
- Two options for large message's offset
 - The offset of the first segment
 - The offset of the last segment



The offset of a large message

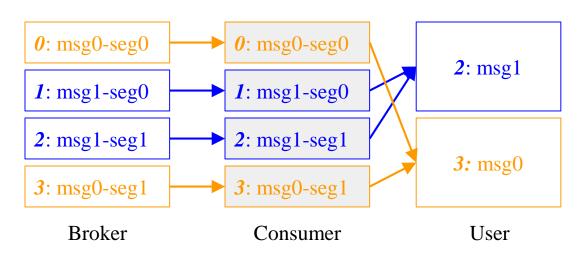
- offset of a large message = offset of first segment
 - First seen first serve
 - Expensive for in-order delivery



Max number of segments to buffer: 4

The offset of a large message

- offset of a large message = offset of last segment
 - Less memory consumption
 - Better tolerance for partially sent large messages.
 - Hard to seek()



Max number of segments to buffer: 3



More details

- Offset Tracking
- Rebalance and duplicates handling
- Producer callback
- Memory management
- Performance overhead
- Compatibility with existing messages
- http://www.slideshare.net/JiangjieQin/handlelarge-messages-in-apache-kafka-58692297



Open Source Clients Library

- Our client library will be open sourced shortly
 - Large message support
 - Auditing



Agenda

- Introduction to Apache Kafka
- Kafka based replication in Espresso
- Message Integrity guarantees
- Performance
- Large message handling
- Security



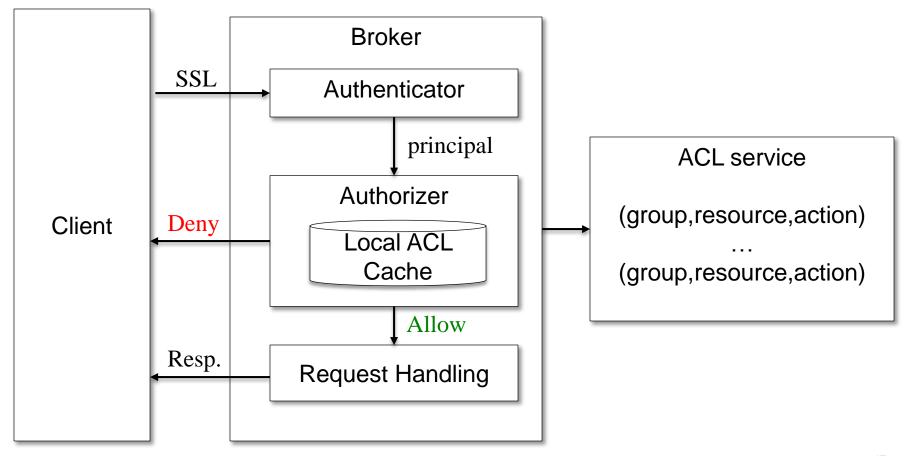
Security in Kafka

- Authentication (SSL, Kerberos, SASL)
- Authorization (Unix-like permission)
 - Resource: Cluster, Topic, Group
 - Operation: Read, Write, Create, Delete, Alter, Describe, Cluster Operation, All
- TLS encryption



Our solution

Authorizer performance is important



Q & A

